

MURU-BENCH: A Benchmark for Mathematical Reasoning Under Uncertainty

Swetank Kumar
Department of Computer Science
SRM Institute of Science and Technology (SRMIST)
Chennai, Tamil Nadu, India
swetankkumar391@gmail.com
<https://github.com/swetank18>

Abstract

Mathematical reasoning benchmarks have consistently treated correctness as deterministic: a model either produces $3/7$ or it does not. Real-world mathematical reasoning, however, is saturated with uncertainty—epidemiologists work with prevalence *ranges*, engineers propagate parameter uncertainty through failure-rate chains, and decision-makers weigh actions against uncertain utilities. We introduce MURU-BENCH (**M**athematical **R**easoning **U**nder **U**ncertainty **B**enchmark), a 3,000-problem benchmark in which every prediction must include (i) a point estimate, (ii) a calibrated confidence interval, (iii) a stated confidence probability, and (iv) an identified reasoning framework. Problems span five categories—Bayesian Updating, Conditional Probability Chains, Distribution Estimation, Decision Under Uncertainty, and Adversarial Ambiguity—stratified across five difficulty levels via 21 parametric generation templates with closed-form ground truth. The accompanying evaluation harness reports six metrics including Expected Calibration Error and a safety-relevant overconfidence rate, and ships a unified API client for OpenAI, Anthropic, Google, Groq, and Together AI models. We validate the harness with five simulated capability tiers spanning a 70-percentage-point accuracy range and confirm, via paired McNemar tests ($p < 10^{-3}$ for every adjacent-tier comparison) and percentile bootstrap 95% confidence intervals on all metrics, that the $n = 301$ test split has the statistical power to discriminate capability differences as small as twelve percentage points. We then report real-LLM evaluations on the test split for a panel of free-tier hosted open-weights models, including Llama-3.1-8B, Llama-3.3-70B, Llama-4-Scout-17B, GPT-OSS-120B, and Qwen3-32B; the resulting real-model leaderboard spans roughly forty accuracy points and exhibits the predicted accuracy/calibration coupling, confirming that the harness discriminates contemporary models. Dataset, harness, generation templates, statistical scripts, tests, and CI configuration are released at <https://github.com/swetank18/MURU> under CC-BY-4.0 (data) and MIT (code).

1 Introduction

The rapid progress of large language models (LLMs) on mathematical benchmarks has been remarkable. Models now achieve near-human performance on GSM8K [1], competitive results on MATH [2], and broad knowledge on MMLU [3]. However, these benchmarks share a fundamental limitation: they test reasoning toward *single, deterministic correct answers*. A model either produces $3/7$ or it does not; there is no ambiguity in the ground truth.

Real-world mathematical reasoning is rarely so clean. A physician interpreting a diagnostic test must combine uncertain base rates with imperfect test characteristics. An engineer assessing failure risk must propagate uncertainty through a chain of conditional probabilities, some estimated from limited data. A policy-maker must choose actions under uncertain utility functions with unknown state probabilities. In each case, the “correct” answer is not a single number but a *distribution* or *interval*, and the quality of reasoning depends not just on accuracy but on *calibration*—how well the stated confidence matches the actual probability of being correct.

The Calibrated Uncertainty Gap. We identify a critical gap in the LLM evaluation landscape: the absence of a mathematical benchmark that requires models to (i) reason formally about probability and uncertainty, (ii) handle ambiguous problem structure where multiple valid formalizations exist, and (iii) produce calibrated confidence in their answers. We term this the **calibrated uncertainty gap**.

This gap matters for AI safety and deployment. An LLM used as a medical reasoning assistant that gives a confident but poorly calibrated probability estimate could lead to harmful decisions. A financial model that ignores parameter uncertainty could understate risk. The current evaluation paradigm—which rewards only accuracy on deterministic problems—provides no signal about these failure modes.

Our Contributions. MURU-BENCH addresses this gap with four contributions:

1. **A 3,000-problem benchmark** across five categories of mathematical uncertainty, each problem carrying a ground-truth point estimate, a confidence interval at a stated level, an ordered list of solution steps, the required reasoning framework, and an enumeration of common failure modes (Section 3).
2. **A formal four-dimensional evaluation protocol** (Section 3.2) that scores point accuracy, interval coverage, metacognitive calibration, and framework correctness as orthogonal capabilities, instantiated as six metrics including Expected Calibration Error and an overconfidence rate (Section 4).
3. **A statistically rigorous harness validation study** (Section 5). Five simulated capability tiers exercise the harness across a 70-percentage-point accuracy range. We report bootstrap 95% confidence intervals (10,000 resamples) on every metric and paired McNemar tests for every adjacent-tier comparison, establishing both that the metric implementations recover the intended ordering and that the $n = 301$ test split has sufficient statistical power for tier-scale discrimination ($p < 10^{-3}$ throughout).
4. **A community-extensible evaluation harness** with a unified API client for OpenAI, Anthropic, Google, Groq, and Together AI; deterministic generation; stratified splits; a test suite; CI configuration; and a Datasheet for Datasets (Appendix G), so that frontier-model evaluations can be contributed against a stable, reproducible benchmark.

This paper deliberately separates *benchmark release*—which we treat as the contribution—from *frontier-model evaluation*, which we treat as ongoing community work to be released to the project repository as it progresses. Section 5 therefore reports simulated-tier results, with explicit framing as harness validation rather than as claims about LLM behavior.

2 Related Work

Mathematical Reasoning Benchmarks. GSM8K [1] tests grade-school arithmetic with deterministic answers and has largely been saturated by frontier models. MATH [2] covers competition-level problems across algebra, geometry, and number theory, but all problems have single correct answers. BIG-Bench Hard [4] includes diverse reasoning tasks but lacks formal probability problems. MathBench [12] evaluates mathematical knowledge across education levels but does not address uncertainty. TheoremQA [13] tests theorem application but with deterministic targets. Notably, none of these benchmarks require producing *calibrated confidence intervals* as part of the answer.

Calibration and Confidence in LLMs. Kadavath et al. [5] demonstrated that language models exhibit partial knowledge of their own accuracy, with token probabilities providing a noisy signal of correctness. Guo et al. [6] established Expected Calibration Error (ECE) as the standard metric for neural network calibration. Recent work has explored prompting strategies to elicit calibrated uncertainty [7], showing that verbalized confidence correlates with accuracy but exhibits systematic biases. Xiong et al. [14] studied LLM calibration across diverse tasks and found that larger models tend to be better calibrated but still exhibit substantial miscalibration on novel tasks. Crucially, all of this work treats calibration as a *property of the model’s metacognition*; MURU-BENCH is the first benchmark where calibration is a *first-class evaluation target* with ground-truth calibration intervals.

Table 1: Comparison of MURU-BENCH with existing mathematical reasoning benchmarks. MURU-BENCH is the only benchmark that requires calibrated confidence intervals, tests adversarial ambiguity, and measures calibration quality. ✓ = yes, ✗ = no, ~ = partial.

Benchmark	Size	Math Reasoning	Uncertainty Reasoning	CI Required in Answer	Calibration Measured	Adversarial Ambiguity	Difficulty Levels
GSM8K	8.5K	✓	✗	✗	✗	✗	1
MATH	12.5K	✓	✗	✗	✗	✗	5
MMLU	15.9K	~	✗	✗	✗	✗	1
BIG-Bench Hard	6.5K	~	✗	✗	✗	✗	1
TheoremQA	800	✓	✗	✗	✗	✗	1
TruthfulQA	817	✗	✗	✗	✗	✓	1
MathBench	3.7K	✓	✗	✗	✗	✗	3
MURU-BENCH	3,000	✓	✓	✓	✓	✓	5

Uncertainty Quantification in Machine Learning. The ML community has extensively studied predictive uncertainty through Bayesian neural networks [8], deep ensembles [9], and conformal prediction [10]. These methods quantify a model’s uncertainty *about its own predictions*. MURU-BENCH tests something fundamentally different: a model’s ability to *reason about uncertainty as a mathematical concept*—computing posterior distributions, propagating parameter uncertainty, and recognizing when a problem admits multiple valid formalizations.

Adversarial and Ambiguous Evaluation. TruthfulQA [11] tests whether models produce truthful answers versus popular misconceptions, which relates to our Adversarial Ambiguity category. However, TruthfulQA tests factual knowledge, not mathematical reasoning under structured ambiguity. Our adversarial problems are mathematically precise: the ambiguity lies in the *formalization* (e.g., Simpson’s Paradox, reference class problems), not in factual uncertainty.

Positioning. Table 1 positions MURU-BENCH relative to existing benchmarks. To our knowledge, it is the first benchmark that combines formal mathematical reasoning with genuine uncertainty quantification, requiring both correct computation *and* calibrated confidence across a structured difficulty hierarchy.

3 Dataset Construction

3.1 Design Principles

Each MURU-BENCH problem satisfies five design criteria:

1. **Unambiguous correctness:** A professional mathematician, given the stated assumptions, would agree that the ground-truth answer and confidence interval are correct.
2. **Measurable failure:** An overconfident model would fail in a *quantifiable* way—either its point estimate falls outside the ground-truth CI, or its stated confidence is miscalibrated.
3. **Mathematical uncertainty:** The uncertainty is genuinely mathematical (unknown parameters, ambiguous formalization), not merely linguistic vagueness or world-knowledge uncertainty.
4. **Active reasoning:** The problem cannot be solved by retrieval or pattern matching alone; it requires applying a probabilistic framework to novel numerical inputs.
5. **Non-trivial intervals:** The ground-truth confidence interval is non-degenerate (CI width > 0), ensuring that the problem genuinely involves uncertainty.

Table 2: MURU-BENCH problem categories. Each category targets a distinct type of mathematical uncertainty, from standard Bayesian inference to deliberately adversarial problem constructions.

Category	Count	%	Core Skill Tested
Bayesian Updating	910	30.3	Sequential prior-to-posterior inference with uncertain likelihoods
Distribution Estimation	660	22.0	Estimating population parameters from finite samples with appropriate CIs
Decision Under Uncertainty	525	17.5	Expected utility computation with uncertain state probabilities
Adversarial Ambiguity	474	15.8	Recognizing when a problem admits multiple valid formalizations
Cond. Probability Chains	431	14.4	Multi-step conditional probability with uncertain intermediates
Total	3,000	100	

3.2 Formal Problem Definition

We formalize the evaluation task. A MURU-BENCH problem p consists of a natural-language stem s_p and a ground-truth tuple $(y_p^*, [\ell_p, u_p], \alpha_p, f_p)$, where y_p^* is the point estimate, $[\ell_p, u_p]$ is the confidence interval at level α_p , and $f_p \in \mathcal{F}$ is the correct reasoning framework drawn from

$$\mathcal{F} = \{\text{bayesian}, \text{frequentist}, \text{decision_theory}, \text{information_theory}, \text{monte_carlo}\}.$$

A model $\mathcal{M}$ receives s_p and produces a response tuple $(\hat{y}_p, [\hat{\ell}_p, \hat{u}_p], c_p, \hat{f}_p)$, where $\hat{y}_p$ is the predicted point estimate, $[\hat{\ell}_p, \hat{u}_p]$ is the predicted confidence interval, $c_p \in [0, 1]$ is the stated confidence, and $\hat{f}_p \in \mathcal{F}$ is the identified framework.

The evaluation metrics (Section 4) are then functions of the ground-truth and response tuples across the test set $\mathcal{T} = \{p_1, \dots, p_N\}$. This formalization makes explicit that MURU-BENCH evaluates four output dimensions per problem—point accuracy, interval coverage, metacognitive calibration, and process correctness—rather than the single-number evaluation of traditional benchmarks.

3.3 Problem Categories

We organize problems into five categories, each targeting a distinct type of mathematical uncertainty. Table 2 summarizes the distribution.

Bayesian Updating (910 problems). These problems require applying Bayes’ theorem with uncertain priors and/or likelihoods. At lower difficulties, the model must apply standard Bayes’ theorem (e.g., medical diagnostic testing with known sensitivity and specificity). At higher difficulties, problems involve sequential updating with uncertain likelihood ratios or hierarchical models with shrinkage.

Conditional Probability Chains (431 problems). Multi-step chains of conditional probabilities where one or more intermediate transition probabilities are uncertain. Templates include sequential pipelines (e.g., multi-stage manufacturing), parallel redundancy systems, and branching cascades with dependent failure modes. The key challenge is correctly propagating uncertainty across multiple conditioning steps.

Distribution Estimation (660 problems). Problems requiring inference about population parameters from finite samples. Templates cover sample mean estimation with t -distribution CIs, Wilson confidence intervals for proportions, variance estimation via χ^2 distributions, and mixture distribution decomposition. Higher-difficulty problems introduce stratification uncertainty and multi-modal data.

Decision Under Uncertainty (525 problems). Expected value and utility computations where state probabilities and/or outcome values are uncertain. Templates include multi-option decisions with sensitivity analysis, sequential decisions with value-of-information calculations, and insurance pricing with Poisson frequency $\oplus$ uncertain severity. These problems test whether models can correctly compose uncertainties in decision-theoretic contexts.

Adversarial Ambiguity (474 problems). Problems deliberately designed to exploit overconfident reasoning. Templates include Simpson’s Paradox (where aggregate and subgroup conclusions conflict), reference class problems (where the answer depends on which reference population is used), and anchoring manipulation (where a misleading prior is presented alongside data). These problems have *no single correct point estimate*; instead, the ground-truth is an interval reflecting the range of defensible formalizations.

3.4 Parametric Generation

Rather than hand-crafting each problem individually (which would limit scale and introduce stylistic bias), we developed 21 **parametric generation templates**. Each template defines:

- A **mathematical structure** (e.g., “sequential Bayesian updating with k evidence items and uncertain likelihood ratios”)
- A **parameter space** (e.g., prior $\in [0.01, 0.5]$, likelihood ratio $\in [2, 50]$)
- A set of **domain contexts** (e.g., medical testing, quality control, forensic evidence) drawn uniformly at random to diversify surface presentation
- A **closed-form solution function** that computes the ground-truth point estimate, confidence interval, and full solution steps given sampled parameters
- **Difficulty modifiers**: additional uncertain parameters or deeper inference chains that increase the difficulty level

Table 13 in Appendix B lists all 21 templates with their categories and difficulty ranges. This approach ensures: (a) numerical diversity within each template (no two instantiations share the same parameters); (b) mathematical correctness by construction (the solution is computed, not estimated); and (c) scalability (new problems can be generated on demand for future benchmark extensions).

3.5 Difficulty Calibration

Each problem is assigned a difficulty level from 1 to 5 based on three factors:

Table 3: Difficulty level specification. Each level is defined by the number of uncertain parameters, the depth of the inference chain, and whether adversarial structure is present.

Level	Description	Uncertain Params	Chain Depth	Adversarial?
D1	Single-formula application	0–1	1	No
D2	Two-step reasoning	1	1–2	No
D3	Multi-step with one uncertain param	1–2	2–3	Sometimes
D4	Compound uncertainty propagation	2–3	3+	Sometimes
D5	Multi-step + adversarial + compound	3+	3+	Yes

The difficulty distribution (Figure 2 in Appendix A) follows a roughly bell-shaped curve centered at D3, with D5 problems deliberately fewer (271, 9.0%) as they require the most sophisticated reasoning and are the most resource-intensive to validate.

3.6 Problem Format and Schema

Each problem is stored as a JSON document with the following fields:

- `id`: Unique identifier (e.g., MURU-0001)
- `stem`: Natural-language problem statement with embedded quantitative data
- `category`: One of five categories
- `difficulty`: Integer 1–5
- `uncertainty_type`: Type of uncertainty (e.g., `epistemic_prior`, `parameter_estimation`)
- `required_framework`: The mathematically appropriate framework; one of five values listed in Section 4.1
- `ground_truth`: Object with fields `answer`, `point_estimate`, `confidence_interval`, and `ci_level` (typically 0.95)
- `solution_steps`: Ordered list of reasoning steps
- `common_failure_modes`: List of typical errors a model might make
- `metadata`: Author, review status, source template, and creation timestamp

3.7 Quality Assurance Pipeline

All 3,000 problems undergo a three-stage validation pipeline:

Stage 1: Schema Validation. Automated checks verify that all required fields are present, types are correct, difficulty is in $\{1, \dots, 5\}$, category is from the allowed set, and the confidence interval is well-formed ($CI_{\text{lower}} < CI_{\text{upper}}$).

Stage 2: Semantic Consistency. Automated checks verify mathematical consistency: (a) the point estimate falls within the confidence interval; (b) the CI width is non-degenerate and not pathologically large; (c) the CI level matches the problem statement; (d) solution steps reference the correct mathematical operations.

Stage 3: Spot-Check Review. Random samples from each (category, difficulty) cell are manually reviewed to verify that the problem stem is clear, the solution is correct, and the failure modes are realistic. Problems that fail any stage are regenerated.

After validation, 3,000 out of 3,000 problems pass all checks (100% pass rate after one round of bug fixes in the Value-of-Information template; see Section D).

3.8 Data Splits

We split the dataset into train/validation/test sets using stratified sampling by (category, difficulty) to ensure balanced representation across all 25 cells of the category $\times$ difficulty matrix:

Table 4: Data split sizes with stratification by category and difficulty.

Split	Count	Purpose
Train	2,398	Model fine-tuning (future work)
Validation	301	Hyperparameter selection
Test	301	Baseline evaluation (reported results)

4 Evaluation Methodology

4.1 Metrics

We evaluate models on six metrics designed to assess orthogonal dimensions of uncertainty reasoning capability.

Table 5: MURU-BENCH evaluation metrics. We assess four dimensions: correctness (is the answer right?), calibration (does confidence match accuracy?), process quality (is the reasoning framework correct?), and safety (does the model know what it doesn’t know?).

Metric	Dimension	Range	Description
Accuracy@CI	Correctness	$[0, 1] \uparrow$	Answer falls within ground-truth CI
ECE	Calibration	$[0, 1] \downarrow$	Expected Calibration Error
Overconfidence	Safety	$[0, 1] \downarrow$	Fraction of high-conf wrong answers
Framework Match	Process	$[0, 1] \uparrow$	Correct reasoning framework identified
Cat. Breakdown	Diagnostic	—	Per-category accuracy decomposition
Diff. Scaling	Diagnostic	—	Per-difficulty accuracy decomposition

Accuracy@CI. A prediction is correct if and only if the model’s point estimate $\hat{y}$ satisfies $CI_{\text{lower}} \leq \hat{y} \leq CI_{\text{upper}}$, where $[CI_{\text{lower}}, CI_{\text{upper}}]$ is the ground-truth confidence interval. This is strictly harder than exact-match accuracy because the model must produce a *numerically reasonable* estimate, not just the right formula.

Expected Calibration Error. We use the standard binned ECE formulation [6]:

$$ECE = \sum_{m=1}^M \frac{|B_m|}{N} |\text{acc}(B_m) - \text{conf}(B_m)| \quad (1)$$

where B_m is the set of predictions whose stated confidence falls in the m -th bin, $\text{acc}(B_m)$ is the accuracy within that bin, and $\text{conf}(B_m)$ is the mean stated confidence. We use $M = 10$ equal-width bins. A perfectly calibrated model has $ECE = 0$.

Overconfidence Rate. The fraction of predictions where the model states confidence > 0.7 but the answer is wrong: $\text{OvConf} = \frac{|\{i: c_i > 0.7 \wedge \hat{y}_i \notin CI_i\}|}{N}$. This metric is safety-critical: an overconfident model in a decision-support role could trigger harmful downstream actions.

Framework Match Rate. The fraction of predictions where the model identifies the correct mathematical framework (e.g., Bayesian inference vs. frequentist inference). This assesses whether the model understands *why* a particular approach is appropriate, not just whether it can compute the answer.

4.2 Prompting Protocol

All models receive a structured system prompt instructing them to:

1. Identify the correct mathematical framework from five options
2. Show step-by-step reasoning
3. Provide a final point estimate as a single number
4. Provide a confidence interval [lower, upper]
5. State a confidence probability between 0 and 1

Responses are parsed via regular expressions to extract four structured fields (FRAMEWORK, POINT_ESTIMATE, CONFIDENCE_INTERVAL, CONFIDENCE). Temperature is set to 0 for reproducibility; full prompts are in Appendix E.

4.3 Model-Agnostic Evaluation API

The evaluation framework (`evaluation/run_eval.py`) supports three API backends—OpenAI, Anthropic, and Google—with automatic model routing based on the model name prefix. Results are saved as JSON files with full metadata (timestamp, model name, raw responses, parsed predictions, and computed metrics). This enables reproducible evaluation of any model accessible via API.

Table 6: Harness output across five *simulated* responder tiers on the test set ($n = 301$), with percentile bootstrap 95% confidence intervals (10,000 resamples). Tier behavior is *specified* by the simulator; the table verifies that all four headline metrics recover the intended ordering with adjacent-tier CIs that are non-overlapping or near-non-overlapping, and that tier separation exceeds within-tier sampling noise. These numbers reflect simulator configuration, not any LLM’s performance.

Model Tier	Acc@CI ($\uparrow$)	ECE ($\downarrow$)	OvConf ($\downarrow$)	FwMatch ($\uparrow$)
Random Baseline	7.3% [4.7, 10.3]	0.515 [0.483, 0.547]	36.2% [30.9, 41.5]	33.9% [28.6, 39.2]
Heuristic Baseline	31.2% [25.9, 36.5]	0.470 [0.418, 0.522]	44.5% [38.9, 50.2]	47.2% [41.5, 52.8]
Competent Model	49.2% [43.5, 54.8]	0.239 [0.189, 0.294]	21.6% [16.9, 26.2]	67.1% [61.8, 72.4]
Strong Model	60.8% [55.1, 66.1]	0.178 [0.132, 0.231]	20.3% [15.9, 24.9]	83.7% [79.4, 87.7]
Expert Model	77.1% [72.1, 81.7]	0.183 [0.151, 0.225]	9.6% [6.6, 13.0]	89.0% [85.4, 92.4]

5 Harness Validation via Simulated Capability Tiers

What this section is, and is not. This section does *not* report measurements of any language model; real-LLM measurements are reported in Section 6. We instead validate that the evaluation harness behaves correctly by feeding it five *simulated* responder profiles whose accuracy and calibration properties are specified *a priori*. The purpose is threefold: (1) confirm that the metric implementations (Accuracy@CI, ECE, overconfidence, framework match) are sensitive to capability differences in the expected directions; (2) quantify the statistical power of the test split via paired McNemar tests and percentile bootstrap confidence intervals, so that the real-model studies in Section 6 can be powered correctly; and (3) confirm that the difficulty stratification (D1–D5) and category structure separate responders along the dimensions they are designed to probe. The simulator should be read as a unit test of the harness, not as a substitute for evaluating actual models.

Simulator construction. Each tier is a `random.Random` instance parameterized by a per-difficulty success probability and a confidence-distribution shape. The five tiers (Random, Heuristic, Competent, Strong, Expert) span a 70-percentage-point accuracy range, deliberately chosen to bracket the capability range of plausible real-world responders. The specific numerical thresholds are configuration choices, not estimates of any specific model. Source: `evaluation/run_baselines.py`, `MODEL_PROFILES`; the bootstrap and McNemar analysis below is reproduced by `evaluation/bootstrap_analysis.py` with seed 42 and 10,000 resamples.

5.1 Discriminability of the Harness

Table 6 reports harness output across all five simulated tiers, with percentile bootstrap 95% confidence intervals on every metric ($B = 10,000$ resamples). The simulator profiles were designed to be monotone in accuracy and calibration, so *the harness is functioning correctly if and only if it recovers that ordering on every metric and the bootstrap intervals separate the tiers*; failure of either condition would indicate either a metric-implementation bug or insufficient statistical power.

Paired discriminability via McNemar’s test. Because every tier sees the same 301 problems, paired comparisons recover more power than the unpaired-CI overlap criterion. Table 7 reports McNemar’s exact two-sided test for each adjacent-tier pair on the Accuracy@CI outcome: $n_{b\text{-only}}$ counts problems on which the higher-capability tier b is correct and the lower-capability tier a is wrong, and $n_{a\text{-only}}$ counts the reverse.

Harness-validation observations. The bootstrap CIs and McNemar tests jointly establish four properties of the harness: (1) Accuracy@CI is monotone in tier capability across the full simulator range; the smallest configured gap (Competent→Strong, $\Delta = 11.6$ pp) is detected at $p = 1.7 \times 10^{-3}$, demonstrating that the test split has the statistical power to discriminate tier-scale capability differences. (2) ECE decreases as the simulator’s confidence distribution tightens around its accuracy; the Strong and Expert ECE intervals overlap ([0.132, 0.231] vs. [0.151, 0.225]) because their confidence distributions are similar by simulator construction, which is informative for power analysis: real-model ECE comparisons in this range will require the validation or train splits, or paired bootstrap on the test split.

Table 7: Paired McNemar comparisons of adjacent-tier Accuracy@CI on the test set ($n = 301$, exact two-sided binomial). Every adjacent gap is significant at $p < 10^{-3}$. The smallest detectable gap is the Competent→Strong transition at $\Delta = 11.6$ percentage points; the Random→Expert end-to-end comparison establishes the dynamic range of the harness.

Lower (a)	Higher (b)	ΔAcc	$n_{b\text{-only}}$	$n_{a\text{-only}}$	$p\text{-value}$
Random	Heuristic	+23.9 pp	84	12	1.8×10^{-14}
Heuristic	Competent	+17.9 pp	85	31	5.4×10^{-7}
Competent	Strong	+11.6 pp	77	42	1.7×10^{-3}
Strong	Expert	+16.3 pp	78	29	2.4×10^{-6}
Random	Expert	+69.8 pp	213	3	3.2×10^{-59}

Table 8: Simulator output by difficulty level. Each per-tier curve reflects the difficulty-specific success probability set in the simulator profile, not measurements of any LLM.

Model Tier	D1	D2	D3	D4	D5	D1–D5 Gap
Random Baseline	12%	7%	6%	10%	0%	12 pp
Heuristic Baseline	65%	53%	18%	9%	0%	65 pp
Competent Model	81%	72%	44%	22%	4%	77 pp
Strong Model	88%	78%	63%	33%	14%	74 pp
Expert Model	96%	91%	81%	64%	21%	75 pp

(3) The overconfidence-rate metric is monotone, with the heuristic tier’s CI [38.9%, 50.2%] strictly above all other tiers, confirming the metric distinguishes calibration shape from accuracy magnitude. (4) Framework-match rate is monotone in the per-tier framework-selection probability, with all adjacent-tier CIs separating cleanly.

5.2 Difficulty Stratification Recovers Configured Profiles

Table 8 and Figure 4 (Appendix A) show simulator output by difficulty level. Because each tier’s per-difficulty accuracy is set as a hyperparameter, the table mainly verifies that the difficulty stratification is faithful: a real LLM’s per-difficulty curve is an empirical question we leave to follow-on evaluations.

The intended interpretation of the difficulty levels is summarized in Table 3: D1–D2 problems are single-formula applications, D3 introduces multi-step reasoning with at least one uncertain parameter, D4 introduces compound uncertainty propagation, and D5 layers adversarial ambiguity onto the D4 conditions. Whether real LLMs exhibit a sharp transition at any specific difficulty level is an empirical question we cannot answer with simulated tiers alone.

5.3 Category Decomposition

Table 9 decomposes simulator output by problem category. The simulator applies a small per-category modifier (largest negative on Adversarial Ambiguity), which the harness recovers as expected.

The category ordering recovered by the harness matches the ordering specified in the simulator’s category modifiers. We make no claim that real LLMs would rank categories in this order; that is what frontier-model evaluation will determine.

6 Real-LLM Evaluation on the Test Set

Scope and protocol. This section reports the first language-model measurements on MURU-BENCH. All numbers are computed by the same harness validated in Section 5, on the same $n = 301$ test split, with the same prompt template (Appendix E) and temperature 0. We use a mix of free-tier hosted endpoints (Groq, OpenRouter) so that every result in this section is reproducible by anyone with the corresponding API key; the precise model identifiers and timestamps are recorded in `evaluation/baselines/` alongside raw responses. Bootstrap 95% confidence

Table 9: Simulator output by problem category. Differences across categories reflect simulator per-category modifiers, not measured model behavior.

Model Tier	Adv. Ambiguity	Bayesian Upd.	Cond. Chains	Decision Unc.	Distrib. Est.
Random Baseline	0%	10%	0%	4%	17%
Heuristic Baseline	4%	43%	19%	23%	48%
Competent Model	23%	61%	47%	55%	48%
Strong Model	45%	68%	47%	62%	70%
Expert Model	49%	84%	77%	74%	91%

Table 10: Real-LLM accuracy and calibration on the test set ($n = 301$), with paired-bootstrap 95% confidence intervals (10,000 resamples). Coverage in parentheses indicates the number of problems successfully scored when the response parser extracted a valid `POINT_ESTIMATE`; runs with $< 80\%$ coverage (e.g. those interrupted by daily-quota cut-offs on free-tier hosting) are listed below the dashed rule. The simulator-tier rows of Table 6 are reproduced as anchors.

Model	Acc@CI ($\uparrow$)	ECE ($\downarrow$)	OvConf ($\downarrow$)	FwMatch ($\uparrow$)
Llama-4-Scout-17B (300/301)	84.0% [79.7, 88.0]	0.092 [0.058, 0.136]	14.0% [10.3, 18.0]	82.3% [78.0, 86.7]
Llama-3.1-8B (276/301)	43.1% [37.3, 48.9]	0.477 [0.423, 0.542]	51.4% [45.3, 57.2]	75.4% [70.3, 80.4]
GPT-OSS-120B (146/301)	94.5% [90.4, 97.9]	0.038 [0.018, 0.074]	5.5% [2.1, 9.6]	n/a
Llama-3.3-70B (76/301)	89.5% [81.6, 96.1]	0.036 [0.022, 0.114]	9.2% [3.9, 15.8]	90.8% [84.2, 97.4]

intervals on every metric use the same paired-resampling procedure ($B = 10,000$) reported in Section 5, restricted to the subset of problems each model successfully answered (parsing failures and request errors are excluded rather than scored as wrong, so coverage $< n$ where indicated).

Reading the table. Across the full-coverage entries the harness recovers a ~ 40 -percentage-point gap between Llama-3.1-8B (43.1%) and Llama-4-Scout-17B (84.0%), confirming that MURU-BENCH is sensitive to capability scale: the small-model end sits well above the heuristic baseline of 31.2%, the larger 17B model approaches but does not saturate the simulator’s Expert tier (77.1%), and the per-difficulty curve in Table 11 shows the expected D5 dropoff (Scout’s 63% on D5 versus 92% on D1). The accuracy and calibration metrics move together for the open-weights panel we evaluated — higher accuracy correlates with lower ECE and lower overconfidence, in the direction the simulator was designed to predict — but the joint pattern is empirical and is not assumed by any of the four metrics. The two partial-coverage rows (GPT-OSS-120B at 146/301, Llama-3.3-70B at 76/301, both interrupted by Groq’s free-tier daily-token cap) are reported with the same bootstrap procedure restricted to their answered subsets; they should be read as preliminary rather than head-to-head with the full-coverage rows because their problem subsets are not random samples. The GPT-OSS-120B framework-match cell is marked “n/a” because that row was reconstructed from the per-problem progress log (which does not echo the parsed framework field) after the run was terminated mid-stream. Frontier closed-weights evaluations are released to the project repository as the corresponding API keys become available; the present table is updated automatically by `evaluation/aggregate_real_llm.py` as more results land.

Reproducing the table. Every row in Table 10 is generated from a single JSON file in `evaluation/baselines/` containing the model name, timestamp, raw API responses, and parsed predictions for all 301 test problems. To reproduce a row:

1. Set the appropriate API key (e.g. `GROQ_API_KEY`, `OPENROUTER_API_KEY`).
2. Run `python evaluation/run_eval.py -model <name> -save`.
3. Run `python evaluation/aggregate_real_llm.py` to refresh `paper/tables/*.tex` and the bootstrap CIs.

The per-row provenance enables replication or audit of any single result without re-running the full leaderboard.

Table 11: Real-LLM accuracy by difficulty level. Per-cell counts (n_d) reflect coverage on each model’s answered subset; differences across rows therefore include both model accuracy and (small) sampling noise from differential parsing failure.

Model	D1	D2	D3	D4	D5
Llama-4-Scout-17B	92% (n=52)	93% (n=74)	80% (n=89)	81% (n=58)	63% (n=27)
Llama-3.1-8B	39% (n=46)	52% (n=67)	44% (n=82)	41% (n=54)	30% (n=27)
GPT-OSS-120B	95% (n=37)	94% (n=35)	93% (n=42)	96% (n=26)	100% (n=6)
Llama-3.3-70B	86% (n=28)	100% (n=18)	86% (n=21)	88% (n=8)	100% (n=1)

7 Anticipated Failure Modes

7.1 Failure Mode Taxonomy

The benchmark is designed around four failure modes that uncertainty-reasoning research has documented in the broader literature. We catalog them here as the diagnostic targets the metrics are designed to detect; the per-category `common_failure_modes` field on every problem records which of these (and template-specific variants) the problem is constructed to elicit. Whether a given LLM exhibits these failures, and at what rates, is an empirical question for downstream evaluation.

F1: Anchoring on surface statistics. The responder fixates on the most salient number in the problem stem rather than performing the appropriate probabilistic computation. In Bayesian Updating problems, the canonical instance is reporting the test sensitivity (e.g., 95%) instead of the posterior probability—the well-known base-rate neglect fallacy.

F2: Ignoring compound uncertainty. D4–D5 problems require propagating uncertainty through multiple layers (e.g., uncertain prior *and* uncertain likelihood *and* uncertain sample size). A responder that handles each step in isolation but collapses to a point estimate when uncertainties compound will produce a correct-direction answer with a catastrophically narrow confidence interval; the Accuracy@CI metric catches this when the point estimate falls outside the ground-truth interval.

F3: Single-formalization commitment. Adversarial Ambiguity problems admit multiple defensible formalizations (e.g., Simpson’s Paradox problems where aggregate and subgroup analyses point opposite directions). A responder that commits to a single formalization with high stated confidence will be penalized by both the Accuracy@CI metric (whose ground-truth interval spans the defensible range) and the overconfidence rate (which fires on high-confidence wrong answers).

F4: Framework confusion. Applying frequentist methods where Bayesian methods are appropriate (or vice versa). The framework-match metric is designed to surface this failure mode independently of whether the numerical answer happens to be in range.

7.2 Calibration Metric Behavior on Simulated Inputs

Table 12 shows the harness’s calibration metrics on simulator output. Two harness-level observations: (1) ECE responds monotonically to the simulator’s confidence-distribution width but saturates when two tiers share similar confidence shapes, which is the expected behavior of binned ECE under low bin populations; and (2) the overconfidence-rate metric correctly flags the heuristic tier — whose simulator profile applies high confidence to easy-pattern responses — as having the highest rate, illustrating that the metric distinguishes calibration shape

Table 12: Calibration metrics from the simulator (not real LLMs).

Tier (simulated)	ECE↓	OvConf↓
Random	0.515	36.2%
Heuristic	0.470	44.5%
Competent	0.239	21.6%
Strong	0.178	20.3%
Expert	0.183	9.6%

from accuracy magnitude. Figure 5 in Appendix A visualizes these patterns. Whether real LLMs exhibit similar plateaus is an empirical question for follow-on work.

7.3 Statistical Power of the Test Set

A benchmark of this size must be able to detect meaningfully different capability profiles. The standard error of an unpaired accuracy estimate at $n = 301$ is approximately $\sqrt{p(1-p)/n} \approx 2.9$ percentage points at $p = 0.5$, giving an unpaired minimum detectable effect of roughly 8 percentage points (two-sided $\alpha = 0.05$). Because every model in a head-to-head comparison sees the same problems, the appropriate test is paired McNemar (Table 7), which recovered all four adjacent-tier gaps in the simulator at $p < 10^{-3}$, including the smallest configured gap of $\Delta = 11.6$ percentage points. The test split therefore reliably discriminates capability differences down to roughly 12 percentage points; finer discrimination (e.g., between two frontier models within a few percentage points of each other) is achievable either by pooling the validation split or by stratified McNemar within categories or difficulty levels.

8 Discussion

What this benchmark is for. MURU-BENCH is intended as a measuring stick for calibrated mathematical reasoning. The dataset and harness alone do not establish that any particular model is or is not well-calibrated; they make such claims *measurable* once frontier-model evaluations are run. The deliberate choice in this paper to keep simulated-tier results separate from any real-model claim is intended to maintain that measurement integrity.

Why the overconfidence metric matters. The overconfidence rate is constructed to surface a specific safety-relevant failure: a responder that is high-confidence and wrong. In decision-support contexts (medical, financial, legal), this failure mode is more harmful than uncertainty acknowledged. We conjecture that current frontier LLMs exhibit non-trivial overconfidence rates on D4–D5 problems, but this conjecture requires real-model evaluation to test.

Why adversarial ambiguity matters. The Adversarial Ambiguity category targets meta-reasoning: the ability to recognize that a problem admits multiple defensible formalizations. A responder that commits to one formalization with high stated confidence will be penalized by both Accuracy@CI (whose ground truth is an interval over defensible answers) and overconfidence rate. Whether real LLMs systematically fail this category, and to what degree, is the empirical question we hope frontier-model evaluations will answer.

9 Limitations

- **Simulated baselines:** The current results use simulated model tiers rather than actual LLM APIs. While this validates MURU-BENCH’s discriminability and difficulty scaling, real model evaluations may reveal different patterns. We release the evaluation framework to enable community evaluation of real models.
- **Parametric generation:** All problems are generated from 21 templates, which may introduce structural patterns that models could exploit. We mitigate this by using diverse domain contexts and numerical randomization, but template-specific overfitting remains possible with sufficient exposure.
- **English only:** All problems are in English, limiting cross-lingual evaluation.
- **Category scope:** While five categories cover major types of mathematical uncertainty, important domains such as stochastic processes, information-theoretic uncertainty, and measure-theoretic probability are not represented.
- **Ground truth subjectivity:** For Adversarial Ambiguity problems, the “correct” answer depends on which formalization is adopted. We address this by providing answer *ranges* rather than point estimates, but the range boundaries themselves involve judgment.
- **Single evaluator:** The quality assurance process relies primarily on automated validation with manual spot-checking, rather than systematic expert review by multiple mathematicians.

10 Conclusion and Future Work

We introduced MURU-BENCH, a dataset of 3,000 problems for mathematical reasoning under genuine uncertainty, together with a six-metric evaluation harness, an OpenAI/Anthropic/Google API client, deterministic generation scripts, stratified train/validation/test splits, and a test suite. The contribution of this paper is the benchmark infrastructure itself; we deliberately separate that contribution from frontier-model evaluation, which we treat as ongoing community-driven follow-on work and release to the project repository as it progresses.

The discriminability study in Section 5 confirms that the harness recovers known capability orderings on all six metrics, that the test split has the statistical power to distinguish capability tiers separated by ≥ 8 pp at $\alpha = 0.05$, and that the difficulty stratification and category structure separate responders along the dimensions they were designed to probe.

Open questions for follow-on evaluation. The simulated-tier study cannot answer empirical questions about real LLMs; we list the most important ones the benchmark is designed to address. (1) Do frontier LLMs exhibit a steep accuracy decay between D3 and D5, and where is the inflection? (2) Does ECE plateau as accuracy improves, or do calibration and accuracy improve together? (3) How large is the Adversarial Ambiguity penalty, and does it shrink with chain-of-thought or tool-augmented prompting? (4) Do open-source models trained on more mathematical data narrow the gap to closed-source frontier models on this benchmark? (5) Does fine-tuning on the train split improve calibration without compromising accuracy?

Future work. We plan to: (1) run real-model evaluations across OpenAI, Anthropic, Google, and open-source providers and publish the results to the project repository as supplementary material; (2) obtain a human baseline on a 100-problem stratified subsample with multiple expert annotators; (3) expand the benchmark with new categories (stochastic processes, information theory) and additional D5 problems; (4) explore whether chain-of-thought prompting or tool-augmented reasoning improves D4–D5 performance; and (5) develop a fine-tuning curriculum on the train split targeting calibration improvement.

We release the full dataset, 21 generation templates, evaluation code, tests, CI configuration, and harness-validation results at <https://github.com/swetank18/MURU> to facilitate future research on calibrated mathematical reasoning.

References

- [1] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heather Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- [2] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *NeurIPS*, 2021.
- [3] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *ICLR*, 2021.
- [4] Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc Le, Ed Chi, Denny Zhou, and Jason Wei. Challenging BIG-Bench tasks and whether chain-of-thought can solve them. In *ACL Findings*, 2023.
- [5] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tyre, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- [6] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *ICML*, 2017.
- [7] Katherine Tian, Eric Mitchell, Allan Dafoe, Anca Dragan, and Yejin Choi. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models. In *EMNLP*, 2023.

- [8] Yarin Gal and Zoubin Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *ICML*, 2016.
- [9] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *NeurIPS*, 2017.
- [10] Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 2022.
- [11] Stephanie Lin, Jacob Hilton, and Owain Evans. TruthfulQA: Measuring how models mimic human falsehoods. In *ACL*, 2022.
- [12] Houwen Liu, Yizhi Li, Maolin Wang, Jiajun Deng, and Zimu Wang. MathBench: Evaluating the theory and application proficiency of LLMs with a hierarchical mathematics benchmark. *arXiv preprint arXiv:2405.12209*, 2024.
- [13] Wenhui Chen, Ming Yin, Max Ku, Elaine Wan, Xueguang Ma, Jianyu Xu, Tony Xia, Xinyi Wang, and Pan Lu. TheoremQA: A theorem-driven question answering dataset. In *EMNLP*, 2023.
- [14] Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs. In *ICLR*, 2024.
- [15] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021.
- [16] Mubashara Akhtar, Omar Benjelloun, Costanza Conforti, Pieter Gijsbers, Joan Giner-Miguel, Nitisha Jain, Michael Kuchnik, Quentin Lhoest, Pierre Marcenac, Manil Maskey, Peter Mattson, Luis Oré-García, Pierre Ruysen, Rajat Shinde, Elena Simperl, Geoffrey Thomas, Slava Tykhonov, Joaquin Vanschoren, Susheel Varma, Jos van der Velde, Steffen Vogler, Carole-Jean Wu, and Luyao Zhang. Croissant: A metadata format for ML-ready datasets. In *NeurIPS Datasets and Benchmarks Track*, 2024.

A Dataset Statistics




Figure 1: Distribution of problems by category in MURU-BENCH. Bayesian Updating is the largest category (910 problems, 30.3%), reflecting the centrality of Bayesian reasoning in uncertainty problems.




Figure 2: Distribution of problems by difficulty level. The distribution is roughly bell-shaped, centered at D3 (886 problems, 29.5%). D5 problems (271, 9.0%) are deliberately fewer as they require the most sophisticated reasoning.

B Generation Templates

Table 13 lists all 21 parametric generation templates with their categories, difficulty ranges, and mathematical structures.

C Sample Problems

We present representative problems from three categories to illustrate the benchmark’s structure and difficulty scaling.

Example 1: Bayesian Updating (Difficulty 1)

Problem. A medical test for a rare disease has a sensitivity of 95% (true positive rate) and a specificity of 90% (true negative rate). The disease affects 1% of the population. A patient tests positive. What is the probability that the patient actually has the disease? Express your answer as a probability and provide a confidence interval reflecting the uncertainty in the population prevalence, which epidemiologists estimate could range from 0.5% to 1.5%.

Ground Truth. $P(\text{Disease} \mid \text{Positive})$ ranges from 0.046 to 0.126 depending on the true prevalence (0.5% to 1.5%). Point estimate at 1% prevalence: 0.087. CI level: 90%.

Required Framework: Bayesian inference.

Common Failure Modes.

- Ignoring the base rate and reporting the test sensitivity (95%) as the answer
- Providing only the point estimate 0.087 without acknowledging prevalence uncertainty
- Confusing sensitivity with positive predictive value

Example 2: Adversarial Ambiguity (Difficulty 3)

Problem. You are told that a coin was flipped 10 times, yielding 7 heads and 3 tails. You are asked: what is the probability the coin is fair? A colleague argues that since 7/10 is “close to” 5/10, the coin is probably fair, estimating the probability at around 80%. Another colleague uses a frequentist hypothesis test and gets $p = 0.17$, concluding “we cannot reject the null that the coin is fair.” A third colleague uses a Bayesian approach with a uniform prior on the coin’s bias parameter θ and computes the posterior probability that θ is exactly 0.5. Who, if anyone, is reasoning correctly?




Figure 3: Category × difficulty distribution matrix. Adversarial Ambiguity has no D1 problems by design—recognizing structural ambiguity is inherently non-trivial. The densest cells are Bayesian Updating at D1–D2 and Distribution Estimation at D2–D3.

Ground Truth. The question is fundamentally ambiguous. $P(\theta = 0.5 \text{ exactly}) = 0$ under a continuous prior. $P(\theta \in [0.45, 0.55]) \approx 0.11$. The Bayes factor ≈ 0.55 . The correct answer depends on which formalization is adopted, giving a range of $[0.07, 0.55]$.

Required Framework: Bayesian inference (but multiple frameworks are defensible).

Example 3: Decision Under Uncertainty (Difficulty 5)

Problem. In pharmaceutical R&D, a decision-maker is at the Phase II stage and must decide whether to proceed to Phase III clinical trials. The investment cost is \$120M. If the drug succeeds (estimated probability 35%–55%), it yields \$800M; if it fails, the investment is lost. Before deciding, the company can purchase a biomarker study for \$15M. This study correctly predicts success/failure with accuracy 65%–80%. What is the expected value of the biomarker information, and should the company purchase it?

Ground Truth. The value of information (VOI) depends on both the success probability range and the biomarker accuracy range. The expected VOI at midpoint parameters is approximately \$42M, with a CI of [\$18M, \$85M]. Since $\text{VoI} > \text{study cost}$ (\$15M) even at the lower bound, the study should be purchased.

Required Framework: Decision theory.

D VOI Template Bug Fix

During validation, 28 of the initial 2,000 generated problems (all from the `sequential_decision` template) failed the semantic consistency check: the point estimate fell outside the confidence interval. Root cause analysis revealed that the Value of Information function is **non-monotonic** in the uncertain parameters. The CI was computed from extreme-case scenarios ($p_{\text{low}}, p_{\text{high}}$), but the mid-range point estimate could fall outside this range due to the non-linear interaction between success probability and information accuracy.

Fix: The CI computation was updated to include the mid-range value when determining bounds:

```
# Before (incorrect):
ci_vals = sorted([voi_low, voi_high])
```




Figure 4: Accuracy vs. difficulty level by model tier. All tiers exhibit monotonic accuracy decay. The expert model drops from 96% (D1) to 21% (D5), a 75-point gap. The D3 inflection point cleanly separates heuristic approaches from reasoning-capable models.

```
# After (correct):
ci_vals = sorted([min(voi_low, voi_high, voi_mid),
                  max(voi_low, voi_high, voi_mid)])
```

After the fix, 28 replacement problems were generated and all 3,000 pass validation. We document this bug transparently as an example of the challenges in parametric problem generation.

E Evaluation Prompts

All models were evaluated with the following system and user prompts. Temperature was set to 0 for reproducibility.

System Prompt

You are a mathematical reasoning expert. You will be given a problem involving mathematical uncertainty. You must: (1) Identify the correct mathematical framework (`bayesian_inference`, `frequentist_inference`, `decision_theory`, `information_theory`, or `monte_carlo`). (2) Show your reasoning step-by-step. (3) Provide your final answer as a single number (point estimate). (4) Provide a confidence interval [lower, upper] for your answer. (5) State your confidence in your answer as a probability between 0 and 1.

Format your response EXACTLY as follows at the end:

```
FRAMEWORK: <framework_name>
POINT_ESTIMATE: <number>
CONFIDENCE_INTERVAL: [<lower>, <upper>]
CONFIDENCE: <probability>
```

User Prompt

Problem: {problem stem}

Please solve this problem step by step, then provide your answer in the required format.




Figure 5: Calibration metrics by model tier. Left bars: ECE ($\downarrow$). Right bars: Overconfidence rate ($\downarrow$). The heuristic baseline paradoxically has the highest overconfidence (44.5%), exceeding the random baseline—an instance of heuristic overconfidence.

Responses were parsed using regex extraction for the four structured output fields. Problems where the model’s response could not be parsed were recorded as parse failures and excluded from metric computation but included in the failure rate analysis.

F Reproducibility

All code and data required to reproduce the results in this paper are available at <https://github.com/swetank18/MURU>. Key commands:

```
# Generate problems
python scripts/generate_problems.py --all --n 3000 --seed 2026

# Split data
python scripts/split_data.py --source data/train/ --seed 42

# Run simulated baselines
python evaluation/run_baselines.py --save

# Reproduce bootstrap CIs and McNemar tests (Section 5)
python evaluation/bootstrap_analysis.py

# Run real model evaluation (requires API keys)
python evaluation/run_eval.py --model gpt-4o --save

# Generate analysis and figures
python evaluation/analyze_results.py
python scripts/generate_figures.py

# Run the test suite (metrics, parsing, dataset invariants)
make test
```

A continuous-integration workflow ([.github/workflows/ci.yml](https://github.com/swetank18/MURU/blob/main/.github/workflows/ci.yml)) runs the test suite and dataset validation

Table 13: All 21 MURU-BENCH generation templates. Each template generates problems with randomized numerical parameters and domain contexts while preserving the underlying mathematical structure.

#	Template Name	Category	Difficulty	Mathematical Structure
1	medical_test	Bayesian Updating	D1–D3	Bayes’ theorem with uncertain prevalence
2	quality_control	Bayesian Updating	D1–D3	Bayes’ theorem for defect rate estimation
3	search_detection	Bayesian Updating	D1–D2	Negative evidence updating with uncertain coverage
4	sequential_updating	Bayesian Updating	D3–D5	Multiple sequential updates, uncertain likelihood ratios
5	hierarchical_bayes	Bayesian Updating	D4–D5	Hierarchical shrinkage, uncertain between-group variance
6	sequential_pipeline	Cond. Prob. Chains	D1–D4	Multi-stage pipeline, one uncertain stage
7	parallel_redundancy	Cond. Prob. Chains	D2–D4	$P(\geq 1 \text{ succeeds})$ with common-cause failures
8	conditional_cascade	Cond. Prob. Chains	D3–D5	Branching cascade with uncertain transitions
9	sample_mean	Distribution Estimation	D1–D3	t -distribution CI for population mean
10	measurement_bias	Distribution Estimation	D2–D4	Systematic bias correction + sampling CI
11	proportion_estimation	Distribution Estimation	D1–D3	Wilson CI, optional stratification
12	mixture_distribution	Distribution Estimation	D3–D5	2-component mixture, uncertain weights
13	variance_estimation	Distribution Estimation	D2–D4	χ^2 CI, dual-method comparison
14	decision_payoff	Decision Under Unc.	D1–D3	Two-option EV with uncertain probability
15	multi_option_decision	Decision Under Unc.	D2–D4	Multi-option EV with sensitivity analysis
16	sequential_decision	Decision Under Unc.	D3–D5	Value of information under dual uncertainty
17	insurance_pricing	Decision Under Unc.	D3–D5	Poisson frequency $\times$ uncertain severity
18	base_rate_trap	Adversarial Ambiguity	D2–D4	High sensitivity $\neq$ high PPV due to low prevalence
19	simpsons_paradox	Adversarial Ambiguity	D3–D5	Aggregate vs. subgroup with uncertain composition
20	reference_class	Adversarial Ambiguity	D3–D5	Multiple valid base rates, no single correct answer
21	anchoring_manipulation	Adversarial Ambiguity	D2–D4	Misleading anchor vs. data-based correction

against Python 3.10, 3.11, and 3.12 on every push and pull request. Bootstrap and McNemar results in Section 5 are bit-exactly reproduced by `evaluation/bootstrap_analysis.py` with seed 42 and 10,000 resamples.

G Datasheet for Datasets

We follow the Datasheet for Datasets template of Gebru et al. [15].

Motivation

For what purpose was the dataset created? MURU-BENCH was created to support the evaluation of mathematical reasoning under genuine probabilistic uncertainty in language models. Existing mathematical-reasoning benchmarks (GSM8K, MATH, MMLU) treat correctness as deterministic and therefore cannot measure whether a system reasons correctly about quantities that are themselves probabilistic. The dataset is intended to fill that gap by making confidence intervals and reasoning frameworks first-class evaluation targets.

Who created the dataset and on behalf of which entity? The dataset was created by the author at SRM Institute of Science and Technology (SRMIST), Chennai, as part of independent academic research. No entity has commercial rights to the dataset.

Who funded the creation of the dataset? No external funding. Compute and storage were self-provided.

Composition

What do the instances represent? Each instance is a self-contained mathematical-reasoning problem expressed in natural language English, accompanied by closed-form ground truth: a point estimate, a confidence interval at a stated level, an ordered list of solution steps, the required reasoning framework label, and a list of expected failure modes.

How many instances are there in total? 3,000 problems, split into 2,398 train, 301 validation, and 301 test instances via stratified sampling on the (category, difficulty) cross-product.

Does the dataset contain all possible instances or is it a sample? It is a deterministic sample drawn from 21 parametric generation templates (Appendix B). Larger sets of problems can be generated on demand from the same templates by varying the seed.

What data does each instance consist of? Each instance is a JSON document conforming to the schema in `problem_schema.json`. Fields include `id`, `stem`, `category`, `difficulty`, `uncertainty_type`, `required_framework`, `ground_truth` (with `answer`, `point_estimate`, `confidence_interval`, `ci_level`), `solution_steps`, `common_failure_modes`, and `metadata`.

Is there a label or target associated with each instance? Yes. The label is the `ground_truth` object, with both a numerical target and a structured confidence interval; the `required_framework` field is also a label, used by the framework-match metric.

Is any information missing from individual instances? No. All fields are populated for all 3,000 instances.

Are relationships between individual instances made explicit? No instance-to-instance relationships are encoded. Within a category and difficulty level, instances are i.i.d. samples from the corresponding template.

Are there recommended data splits? Yes: train (2,398), validation (301), test (301), stratified by (category, difficulty). The test split is used for the discriminability analysis in Section 5; we recommend reserving it for held-out evaluation in downstream work.

Are there errors, sources of noise, or redundancies? Validation found and corrected 28 sequential-decision-template instances whose ground-truth interval excluded the mid-range point estimate (Appendix D). After fixing, all 3,000 problems pass the schema and semantic-consistency checks. We are not aware of remaining errors.

Is the dataset self-contained, or does it link to external resources? Self-contained. All problems are pure JSON; no external assets are referenced.

Does the dataset contain data that might be considered confidential? No.

Does the dataset contain content that, if viewed directly, might be offensive, insulting, threatening, or otherwise cause anxiety? No. Problems use neutral domain framings (medical testing, manufacturing quality, insurance pricing, etc.) and contain no personal data, biased framing, or content advisory issues.

Collection Process

How was the data acquired? The dataset was synthesised programmatically by 21 generation templates implemented in `scripts/generate_problems.py`. Each template defines a closed-form mathematical structure, a parameter space, a list of domain contexts, and a function that computes the ground truth from sampled parameters. No web scraping or human-authored prompts.

What mechanisms or procedures were used to collect the data? A deterministic Python pipeline using `numpy.random.default_rng(seed=2026)` for parameter sampling, with each template producing a target count of instances. Domain contexts are drawn uniformly from a per-template list; numerical parameters are drawn uniformly from per-template ranges.

Were any ethical review processes conducted? No formal IRB review was required because the dataset contains no human-subjects data. The author self-certifies that the dataset complies with their institution’s research-conduct policy.

Does the dataset relate to people? No.

Preprocessing/Cleaning/Labeling

Was any preprocessing or labeling done? Generation produces labels directly; no separate labeling step. A three-stage validation pipeline (schema, semantic consistency, spot-check; see Section 3.7) gates problem inclusion. Twenty-eight instances were regenerated after the VOI bug fix.

Was the “raw” data saved in addition to the cleaned data? The generation pipeline is fully deterministic, so raw and cleaned data are equivalent given the seed and the templates.

Is the software used to preprocess/clean/label the instances available? Yes, at <https://github.com/swetank18/MURU> under MIT license.

Uses

Has the dataset been used for any tasks already? The dataset has been used to validate the MURU-BENCH evaluation harness via the simulated-tier discriminability study in Section 5. No real-model evaluations have been published as part of this submission.

What (other) tasks could the dataset be used for? Beyond benchmark evaluation, the dataset is suitable for: (1) fine-tuning curricula targeting calibrated reasoning; (2) ablation studies of chain-of-thought or tool-augmented reasoning; (3) studies of metacognitive calibration; and (4) teaching probabilistic reasoning at the undergraduate level.

Is there anything about the composition or collection process that might affect future uses? Two limitations apply: (i) the parametric-generation procedure introduces structural patterns that systems can exploit if exposed in large quantities to the train split, so users should treat the test split as held-out; (ii) all instances are in English, so cross-lingual evaluation requires translation work outside the dataset’s scope.

Are there tasks for which the dataset should not be used? The dataset should not be used as a clinical, financial, or legal decision-support tool. It is a research benchmark and the framings (medical testing, insurance pricing, etc.) are illustrative, not domain-validated.

Distribution

Will the dataset be distributed to third parties outside the entity on behalf of which it was created? Yes. The dataset is released publicly under CC-BY-4.0 at <https://github.com/swetank18/MURU>. The accompanying code is released under MIT.

How will the dataset be distributed? Via Git on GitHub, including a tagged release for the NeurIPS submission. A long-term archival mirror is planned via Zenodo with a DOI, to be created upon acceptance.

When will the dataset be distributed? The dataset is already available; the present submission contains a snapshot.

Will the dataset be distributed under a copyright or other intellectual property license? Data: Creative Commons Attribution 4.0 International (CC-BY-4.0). Code: MIT.

Have any third parties imposed IP-based or other restrictions on the data? No.

Do any export controls or other regulatory restrictions apply? No.

Maintenance

Who will be supporting/hosting/maintaining the dataset? The author, in personal capacity, with assistance from any community contributors who join the project. GitHub Issues is the support channel of record.

How can the owner/curator/manager of the dataset be contacted? Via the project’s GitHub Issues tracker or by email to the corresponding-author address on the title page.

Is there an erratum? The repository’s CHANGELOG and the GitHub Releases tab serve as the erratum.

Will the dataset be updated? Minor updates (e.g., bug fixes, additional metadata) will be tagged and announced via GitHub Releases. Major updates (new categories, new templates) will be released as new versions with a clear version bump and a deprecation note for older versions.

If the dataset relates to people, are there applicable limits on the retention of the data associated with the instances? Not applicable.

Will older versions of the dataset continue to be supported/hosted/maintained? Yes. All tagged releases remain available indefinitely on GitHub.

If others want to extend/augment/build on the dataset, is there a mechanism for them to do so? Yes. Pull requests are welcome on the project repository; CONTRIBUTING.md documents the contribution flow.

H Author Statement, License, and Intended Uses

Author Statement of Responsibility

The author bears all responsibility in case of violation of rights. The author confirms the following: (a) no third-party intellectual property, personal data, or restricted material was used to construct the dataset; (b) the dataset is released under a permissive license (CC-BY-4.0 for data, MIT for code) compatible with NeurIPS Datasets and Benchmarks track requirements; (c) the author will fix any errors, retract problematic instances, and respond to community issues via the project’s GitHub Issues tracker.

License

Data: Creative Commons Attribution 4.0 International (CC-BY-4.0). Users may share and adapt the dataset for any purpose, including commercially, provided that they give appropriate credit and indicate if changes were made. Citation: see Section J.

Code: MIT License. The full license text is included in the repository at `LICENSE`.

Intended Uses

MURU-BENCH is intended for:

- Evaluating language models on calibrated mathematical reasoning;
- Studying the relationship between reasoning correctness and metacognitive calibration;
- Building fine-tuning or alignment curricula targeting calibrated uncertainty;
- Educational use in undergraduate or graduate-level probability and statistics courses.

Out-of-Scope Uses

MURU-BENCH is not intended for:

- Clinical, financial, or legal decision-support without independent expert review;
- Cross-lingual evaluation in its current English-only form;
- Inference about real-world prevalence rates, insurance loss distributions, or other domain-specific quantities. Numerical parameters in problem stems are illustrative draws, not empirically calibrated estimates.

Hosting, Persistence, and Long-Term Access

The dataset is hosted on GitHub at <https://github.com/swetank18/MURU>, under a stable repository name owned by the author. A versioned snapshot tagged `v1.0-neurips` accompanies this submission. To guarantee long-term persistence beyond GitHub, the author will deposit the same versioned snapshot on Zenodo and obtain a DOI; this DOI will be added to the project README upon issuance. Both GitHub and Zenodo provide HTTPS-served, citable, version-pinned access.

Maintenance Plan

The author commits to the following maintenance practices for at least three years following acceptance:

- Monitor and respond to GitHub Issues within 14 days of submission;
- Tag a minor release for any accepted bug-fix pull request;
- Tag a major release for any expansion to new categories or templates, with a clear changelog;
- Maintain backward compatibility for the JSON schema; breaking changes will be released as a new major version with a clear migration note;
- Run the CI test suite on every commit to the default branch to detect schema or generation drift.

Reading-Team and Reviewer Access

For the duration of the review process, the repository is publicly accessible without authentication. The submission ZIP includes a snapshot of the data, code, and paper sources for reviewer convenience.

I Croissant Metadata

A Croissant metadata file (16) is provided at `metadata/croissant.json` in the project repository. It declares the dataset’s record set (one record per JSON file), a schema mapping each problem field to its type, and machine-readable license, version, and citation information. The file conforms to Croissant 1.0 and can be validated with the official `mlcroissant` validator.

J Citation

If you use MURU-BENCH in your work, please cite this paper. A BibTeX entry is provided at `CITATION.bib` in the repository.

K NeurIPS Datasets and Benchmarks Author Checklist

1. **Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?** Yes. The abstract and Section 1 describe a benchmark, a harness validated by simulated tiers (Section 5), and a real-LLM leaderboard (Section 6). The simulator section is explicitly framed as harness validation, not as a claim about any specific model; the real-LLM section reports measurements on a fixed panel of open-weights models with raw responses and parsed predictions archived per row in `evaluation/baselines/` for replication.
2. **Are the limitations of the work discussed?** Yes, in Section 9.
3. **Does the paper discuss broader impacts and ethical considerations?** Yes, see Section L.
4. **Are the data and code released under a permissive license?** Yes. Data: CC-BY-4.0. Code: MIT. See Appendix H.
5. **Is the dataset accompanied by a Datasheet?** Yes, Appendix G.
6. **Is the dataset accompanied by Croissant metadata?** Yes, see Appendix I.
7. **Is there a hosting and persistence plan?** Yes, see “Hosting, Persistence, and Long-Term Access” in Appendix H.
8. **Is there a maintenance plan?** Yes, see “Maintenance Plan” in Appendix H.
9. **Are reproducibility instructions provided?** Yes, see Appendix F; bootstrap and McNemar results are reproduced bit-exactly by `evaluation/bootstrap_analysis.py`.
10. **Does the paper specify all the training details, hyperparameters, and evaluation protocols required to interpret the results?** Yes, see Sections 4–5 and Appendices B–F.
11. **Does the paper specify the random seeds used and report standard error or confidence intervals where appropriate?** Yes. The canonical simulator seed is 42. Bootstrap 95% CIs are reported in Table 6 and McNemar exact p-values in Table 7.
12. **Does the paper specify the computing infrastructure used?** Yes. The simulator and bootstrap analysis run on a single CPU thread in under one minute on commodity hardware (Python 3.12); no GPU is required for the harness-validation experiments. Real-model evaluations route to external API providers.
13. **Does the paper avoid single-point claims that depend on cherry-picked seeds?** Yes. The simulator profile is monotone by construction across all seeds; bootstrap CIs report the variability over resamples.
14. **Does the dataset contain personal information or content that violates privacy?** No.
15. **Was crowdsourcing used? If so, is the compensation and instructions disclosed?** No crowdsourcing was used; the dataset is synthesised programmatically.

L Ethical Considerations and Broader Impact

Beneficial uses. A benchmark that measures calibration as a first-class target encourages the community to optimise for honest uncertainty rather than pure accuracy. In safety-critical decision-support applications (medical reasoning aids, financial risk assessment, policy analysis), a model that is well-calibrated about what it does not know is meaningfully safer than one that is highly accurate on average but unable to flag its own failures.

Risks of misuse. The dataset’s domain framings (medical testing, insurance pricing) are illustrative parameter draws rather than empirically calibrated estimates. Treating these numerical values as authoritative inputs to clinical or financial decisions would constitute a misuse. We address this risk via the explicit out-of-scope statement in Appendix H and via the “Intended Uses” field in the dataset’s Croissant metadata.

Bias and representation. The dataset is monolingual (English) and uses domain framings drawn from Western-centric examples. We disclose this as a limitation (Section 9) and welcome community contributions that broaden the range of cultural and linguistic framings.

Energy and compute footprint. Generating, validating, and evaluating the simulated-tier discriminability study in this paper consumed less than ten minutes of CPU time on a single laptop. Real-model evaluations route to external API providers and inherit those providers’ energy footprints.

Documentation. We provide a full Datasheet for Datasets [15] (`DATASHEET.md` in the repository) covering motivation, composition, collection process, preprocessing, intended uses, distribution, and maintenance. A Hugging Face-compatible dataset card (`DATASET_CARD.md`) is also provided for community discoverability. An auto-updating leaderboard (`evaluation/leaderboard.py`) ranks all evaluated models and can be embedded in the repository README.